

```
In [20]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import FunctionTransformer
import re
import warnings
from scipy.stats import uniform
warnings.filterwarnings('ignore')
```

```
In [21]: # Load the datasets
train_data = pd.read_csv('C:/Users/v/Desktop/MAYURA/Data Science Projects/Twitter S
validation_data = pd.read_csv('C:/Users/v/Desktop/MAYURA/Data Science Projects/Twit

# Rename columns
train_data.columns = ['id', 'topic', 'sentiment', 'text']
validation_data.columns = ['id', 'topic', 'sentiment', 'text']

# Display the first few rows of the datasets
print(train_data.head())
print(validation_data.head())
```

```
      id      topic sentiment \
0  2401  Borderlands  Positive
1  2401  Borderlands  Positive
2  2401  Borderlands  Positive
3  2401  Borderlands  Positive
4  2401  Borderlands  Positive

      text
0  im getting on borderlands and i will murder yo...
1  I am coming to the borders and I will kill you...
2  im getting on borderlands and i will kill you ...
3  im coming on borderlands and i will murder you...
4  im getting on borderlands 2 and i will murder ...

      id      topic  sentiment \
0  3364   Facebook  Irrelevant
1   352     Amazon   Neutral
2  8312  Microsoft  Negative
3  4371     CS-GO   Negative
4  4433     Google   Neutral

      text
0  I mentioned on Facebook that I was struggling ...
1  BBC News - Amazon boss Jeff Bezos rejects clai...
2  @Microsoft Why do I pay for WORD when it funct...
3  CSGO matchmaking is so full of closet hacking,...
4  Now the President is slapping Americans in the...
```

```
In [22]: # Check for missing values and data types
print(train_data.info())
print(train_data.describe())
print(validation_data.info())
print(validation_data.describe())

# Distribution of the target variable
```

```
print(train_data['sentiment'].value_counts())
print(validation_data['sentiment'].value_counts())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74682 entries, 0 to 74681
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           74682 non-null  int64
1   topic        74682 non-null  object
2   sentiment    74682 non-null  object
3   text         73996 non-null  object
dtypes: int64(1), object(3)
memory usage: 2.3+ MB
None
```

```


              id
count  74682.000000
mean    6432.586165
std     3740.427870
min         1.000000
25%     3195.000000
50%     6422.000000
75%     9601.000000
max    13200.000000
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           1000 non-null  int64
1   topic        1000 non-null  object
2   sentiment    1000 non-null  object
3   text         1000 non-null  object
dtypes: int64(1), object(3)
memory usage: 31.4+ KB
None
```


```

              id
count    1000.000000
mean     6432.088000
std      3728.310569
min         6.000000
25%     3247.750000
50%     6550.000000
75%     9661.750000
max    13197.000000
sentiment
Negative      22542
Positive      20832
Neutral       18318
Irrelevant    12990
Name: count, dtype: int64
sentiment
Neutral        285
Positive       277
Negative       266
Irrelevant     172
Name: count, dtype: int64
```

```
In [23]: # Visualizing the distribution of the target variable in the training set
sns.countplot(x='sentiment', data=train_data)
plt.title('Distribution of Sentiment in Training Data')
plt.show()
```



```
In [24]: # Visualizing the distribution of the target variable in the validation set
sns.countplot(x='sentiment', data=validation_data)
plt.title('Distribution of Sentiment in Validation Data')
plt.show()
```



```
In [25]: # Remove any rows with NaN values in the 'text' or 'sentiment' columns
train_data.dropna(subset=['text', 'sentiment'], inplace=True)
validation_data.dropna(subset=['text', 'sentiment'], inplace=True)

# Check for null values again
print(train_data.isnull().sum())
print(validation_data.isnull().sum())
```

```
id          0
topic       0
sentiment   0
text        0
dtype: int64
id          0
topic       0
sentiment   0
text        0
dtype: int64
```

```
In [26]: # Text preprocessing function
def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'^a-z\s', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    return text
```

```
In [27]: # Apply text preprocessing
train_data['text'] = train_data['text'].apply(preprocess_text)
validation_data['text'] = validation_data['text'].apply(preprocess_text)
```

```
In [29]: # Splitting the data into training and validation sets
X_train, X_test, y_train, y_test = train_test_split(train_data['text'], train_data[

# Display the shapes of the splits
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(59196,) (14800,) (59196,) (14800,)
```

```
In [30]: # Initializing TF-IDF Vectorizer
tfidf = TfidfVectorizer(max_features=10000, ngram_range=(1,3))

# Function for Model Evaluation
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    print("Accuracy: ", accuracy_score(y_test, y_pred))
    print("Classification Report:\n", classification_report(y_test, y_pred))
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.title('Confusion Matrix')
    plt.show()
```

```
In [31]: # Building a pipeline for the model
pipeline = Pipeline([
    ('tfidf', tfidf),
    ('logreg', LogisticRegression(solver='liblinear'))
])
```

```
In [32]: # Hyperparameter tuning
param_dist = {
    'logreg__C': uniform(0.01, 10),
```

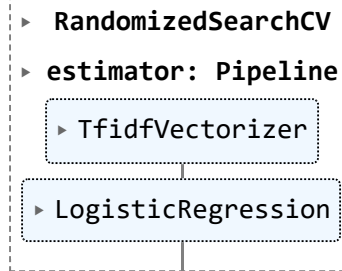
```

    'logreg__penalty': ['l2', 'l1']
}
random_search = RandomizedSearchCV(pipeline, param_distributions=param_dist, n_iter=100)
random_search.fit(X_train, y_train)

```

Fitting 5 folds for each of 50 candidates, totalling 250 fits

Out[32]:



In [33]:

```

# Best parameters
print("Best parameters found: ", random_search.best_params_)

```

Best parameters found: {'logreg__C': 9.498855372533333, 'logreg__penalty': 'l1'}

In [34]:

```

# Best model
best_model = random_search.best_estimator_

```

In [35]:

```

# Model Evaluation on test set
evaluate_model(best_model, X_test, y_test)


# Model Evaluation on validation set
evaluate_model(best_model, validation_data['text'], validation_data['sentiment'])

```

Accuracy: 0.7675675675675676

Classification Report:


	precision	recall	f1-score	support
Irrelevant	0.76	0.68	0.72	2696
Negative	0.81	0.82	0.81	4380
Neutral	0.72	0.74	0.73	3605
Positive	0.76	0.79	0.78	4119
accuracy			0.77	14800
macro avg	0.77	0.76	0.76	14800
weighted avg	0.77	0.77	0.77	14800



Accuracy: 0.913

Classification Report:

	precision	recall	f1-score	support
Irrelevant	0.92	0.90	0.91	172
Negative	0.91	0.93	0.92	266
Neutral	0.93	0.89	0.91	285
Positive	0.89	0.92	0.91	277
accuracy			0.91	1000
macro avg	0.91	0.91	0.91	1000
weighted avg	0.91	0.91	0.91	1000



In []: